

PART I — FOUNDATIONS

Chapter 1

Why RAG? The Problem with Plain LLMs

“If an answer sounds confident, fluent, and detailed — and is completely wrong — the problem is not the user. The problem is the system.”

Chapter 1 — Why RAG? The Problem with Plain LLMs

Before we write a single line of code, we need to answer a question that almost every tutorial skips: why does RAG exist at all? Large Language Models are already astonishingly capable. They can summarize, translate, draft, code, explain, debate, and brainstorm. So why bolt anything onto them? Why not just ask the model and call it a day?

The honest answer is that a plain LLM, taken on its own, is a brilliant graduate who has read an enormous library — but only the editions that existed up to a particular date, has no access to the company intranet, cannot consult the latest journal, and will sometimes invent a citation rather than admit ignorance. For casual chat, that is fine. For a production system that must give correct, current, source-grounded answers about your data, it is not.

This chapter builds that argument carefully. We will look at what an LLM actually is under the hood, where its limitations come from, why the two obvious workarounds — better prompting and fine-tuning — only solve part of the problem, and finally, what unique gap Retrieval-Augmented Generation fills. By the end, you will not just know that RAG is useful; you will know precisely which failure modes of plain LLMs it fixes, and which it does not.

1.1 What Large Language Models actually are

Strip away the marketing, and a Large Language Model is a single mathematical function — astonishingly large, but conceptually simple. It takes a sequence of text as input, and it outputs a probability distribution over what should come next. That is its entire job. Everything else — chat, code generation, summarization, reasoning — is an emergent consequence of doing this one thing extraordinarily well, billions of times, across nearly every kind of text humans have written.

Definition — Large Language Model (LLM)

A neural network, typically based on the Transformer architecture, that has been trained on massive amounts of text to predict the next token in a sequence. By repeatedly predicting one token at a time, it can generate fluent, coherent, and contextually appropriate text across an enormous range of tasks.

Two adjectives in that definition do most of the work. The model is large in two senses: it has many parameters (often tens or hundreds of billions of trainable weights), and it has seen a huge amount of training text (typically measured in trillions of tokens). And it is a model of language — meaning, of statistical patterns in human writing — not of facts, truth, or reality. This distinction matters enormously, and it sits at the heart of every problem we will discuss in this chapter.

Analogy — The voracious reader with no library card

Imagine someone who, by some miracle, has read every book, paper, manual, forum post, and source-code repository printed before a certain date — and absorbed not the facts in them but the patterns of how words follow other words. They can write convincingly in any style, on almost any topic, because

they have an internal sense of what such writing usually sounds like. But once they leave the room, they cannot consult a single new book. Whatever the world publishes after that date is invisible to them. That is a Large Language Model.

This framing — pattern-matcher rather than fact-checker — explains both the power of plain LLMs and almost all of their failure modes. They are spectacular at producing text that looks like a correct answer. Whether the text is, in fact, a correct answer is a separate question that the model itself is structurally unable to verify. That is the problem RAG was designed to fix.

1.2 The Transformer architecture in one page

Almost every modern LLM is built on the Transformer architecture, introduced by Vaswani et al. in 2017. You do not need to derive its equations to use it well, but a working mental model of how it processes language will pay dividends for the rest of this book. Here is the entire architecture, condensed to one page of plain English.

A Transformer takes a sequence of tokens (numerical IDs that represent pieces of text) and runs them through a stack of identical layers. Each layer does two things: it lets every token look at every other token in the sequence using a mechanism called self-attention, and then it transforms each token's representation through a small feed-forward network. After many such layers, every token's vector representation has been refined into something deeply context-aware — the word "bank" in "river bank" ends up with a different internal representation than the word "bank" in "investment bank", purely from the context surrounding it.

Definition — Self-attention

A mechanism that lets each token in a sequence directly attend to (look at, weigh, and incorporate information from) every other token, regardless of how far apart they are. This is what allows a Transformer to capture long-range dependencies — for example, that the pronoun in the last sentence of a paragraph refers to a noun in the first sentence.

There are two flavors of Transformer worth distinguishing. Encoder models (like BERT) read the entire input at once and produce rich representations of it — they are excellent for classification, similarity, and embedding tasks. Decoder models (like the GPT family, Llama, Claude, Gemini, and most chatbots) generate text one token at a time, each new token conditioned on everything that came before. Most of the LLMs we will integrate with in this book are decoder models.




Figure 1.1 — A decoder-only Transformer in one diagram. Each layer (self-attention + feed-forward) is repeated N times to deepen contextual understanding.

The crucial thing to remember is what comes out at the end: not an answer, not a fact, but a probability distribution over the next token. The model does not know it is correct. It only knows what is statistically likely to come next, given everything it has read during training. We will return to this point repeatedly — it is the seed of the hallucination problem.

1.3 Tokens, embeddings, and context windows

Three terms appear constantly in LLM work, and confusion about them blocks understanding of every later topic. Let us pin them down precisely.

Tokens

LLMs do not see characters or whole words. They see tokens — pieces of text, often subword fragments, that the model's tokenizer has learned to produce. Common words become a single token; rare words and arbitrary strings get broken into multiple tokens.

Example — How a tokenizer sees text

Input string: "ChatGPT is amazing"

Tokenized as: ["Chat", "GPT", " is", " amazing"]

Token IDs: [13313, 38, 1041, 8186]

Notice that "ChatGPT" splits into two tokens, while " is" and " amazing" each include the leading space — that is intentional, not a bug.

Why this matters in practice: every API charges by tokens, every model has a maximum number of tokens it can consider at once, and every chunking strategy in our ingestion pipeline (Chapter 6 onward) will be measured in tokens. As a rough rule of thumb for English, one token $\approx$ 0.75 words, or about 4 characters.

Embeddings

Once a token has been converted to an integer ID, the model immediately maps that ID to a dense vector — a list of numbers — called an embedding. The dimensionality is fixed by the model's architecture: a small open-source model like all-MiniLM-L6-v2 produces 384-dimensional vectors; OpenAI's text-embedding-3-large produces 3,072. These vectors are not random. They are positions in a high-dimensional space where geometric closeness corresponds to semantic similarity.

Definition — Embedding

A dense numerical vector that represents the semantic meaning of a piece of text. Texts with similar meanings produce vectors that are close together in this vector space; texts with unrelated meanings produce vectors that are far apart.

Analogy — GPS coordinates of meaning

Think of an embedding as the GPS coordinates a sentence inside a vast "meaning map." "How do I cancel my subscription?" and "How can I unsubscribe from my plan?" land on almost the same street corner, even though they share almost no words. "How do I bake sourdough?" lives on a completely different continent. This is why we can find relevant documents by comparing vector distances rather

than matching keywords — and it is the mathematical foundation of every retrieval system in this book.

Context windows

The context window is the maximum number of tokens an LLM can read and reason over in a single call. It is the model's working memory — anything outside it might as well not exist. Smaller models offer 8K or 32K tokens; larger ones go up to 128K, 200K, or even 1M+ tokens.

A common beginner intuition is: "Context windows are huge now, so I'll just paste my entire 800-page company manual into every prompt, and the model will figure it out." There are three problems with this. First, cost: every token in every call costs money and time. Second, latency: the model's processing time grows with input length. Third, and most subtly, retrieval-from-context degrades for long inputs — research shows that LLMs disproportionately attend to the start and end of very long prompts, often missing crucial information stuffed in the middle. This phenomenon is sometimes called "lost in the middle."

This is precisely the gap RAG fills. Instead of sending an entire library on every call, RAG sends only the few hundred tokens that are actually relevant to the user's current question — drawn from a much larger external store via similarity search.

1.4 The hallucination problem — concrete examples

Hallucination is the most-talked-about failure mode of LLMs, and the most misunderstood. It is not the model being "buggy" or "lying." It is the model doing exactly what it was trained to do — predicting plausible next tokens — in a situation where the truth happens not to be the most plausible-sounding continuation.

Definition — Hallucination

A confident, fluent, and well-formed output from an LLM that contains incorrect, fabricated, or unsupported information. The model is not aware that the output is wrong; from its perspective, it has simply produced a statistically reasonable continuation of the prompt.

Analogy — The overconfident friend

We all know one. Ask them anything — the GDP of Bolivia, the year a film was released, the meaning of an obscure Latin phrase — and they will answer with total conviction. Half the time they are right. The other half they are quietly making it up, and their tone never gives away which half you are listening to. Plain LLMs have exactly this personality. The cure is not to teach them to be quieter; it is to give them a reference book they must check before they speak. That reference book is RAG.

Concrete examples

Example 1 — Fabricated citations. Ask a plain LLM, "What does the 2023 paper by Smith et al. on retrieval-augmented generation say about reranking?" and you may receive a beautifully structured paragraph with quotes, page numbers, and a journal name — for a paper that does not exist. The model has seen thousands of real citations and reproduces the form perfectly; the content is improvised.

Example 2 — Confident arithmetic errors. Ask "What is 8.11 minus 8.9?" and an undertrained or smaller model may confidently respond with 0.21 or even something stranger, because tokens-as-arithmetic is not the same skill as text-as-language. The answer is wrong, but the explanation around it sounds entirely competent.

Example 3 — Plausible but invented company facts. Ask, "What is the warranty period on our XR-440 industrial controller?" and a plain LLM that has never seen your internal documentation will produce an answer anyway — usually "two years" or "one year," because those are the most common warranty periods in its training data. The number sounds reasonable. It happens to be wrong.

Example 4 — Stale facts presented as current. Ask, "Who is the current CEO of company X?" and you may receive the name of someone who left the role two years ago. The model is not lying; it simply has no mechanism to know that the world has moved on since its training data was collected.

Notice the common thread. In every case, the output is fluent, well-formatted, often beautifully structured — and wrong. Hallucination is not a bug we can patch out by writing a stricter prompt. It is a structural consequence of asking a probabilistic next-token predictor to behave like a database. The fix is to add a database.

1.5 Knowledge cutoffs and stale information

Every LLM has a training data cutoff — a date past which it has seen no text. The model could be released yesterday, but if its training corpus stopped collecting in October 2023, it has zero awareness of anything that happened between October 2023 and now. From the model's perspective, the post-cutoff world simply does not exist.

Definition — Knowledge cutoff

The point in time at which an LLM's training data ends. Information about events, facts, products, prices, people, or laws that came into existence (or changed) after the cutoff is unknown to the model and cannot be reasoned about — though, problematically, the model will often try to reason about it anyway.

Two failure modes follow directly from this. The first is silence: ask about a recent event and the model says, "I do not have information past my training cutoff." That is the polite, useful failure mode. The second is hallucination: the model attempts to answer anyway, fabricating plausible-sounding details. Which mode you get depends on the model, the prompt, and luck.

Even within the cutoff, there is a deeper problem: information staleness. Consider a model whose data was collected in 2023. It probably contains a 2019 web page that says "the current CEO of Acme Corp is Jane Doe," a 2021 page saying the same, and a 2023 page announcing her replacement. All three statements are now part of the training data; the model has no concept of "latest." It may quote whichever version dominates its weights, or blend them. There is no version of prompt engineering that fixes this — the underlying data cannot be re-sorted at inference time.

This is the second core gap RAG fills. By keeping the knowledge base outside the model and updating it independently, we decouple "the model's reasoning ability" from "the world's current state." The model stays the same; we re-ingest the new documents.

1.6 Why prompt engineering alone is not enough

Prompt engineering is the practice of carefully designing the input to an LLM — instructions, role assignments, examples, constraints — to coax better outputs. It is genuinely powerful, and the third part of this book is dedicated to it. But it is not a complete solution to the problems we have just described and confusing it with one is one of the most common beginner mistakes.

Here is what prompt engineering can do. It can clarify what you want ("answer in three sentences," "return JSON," "explain the reasoning step by step"). It can supply examples (few-shot prompting). It can assign a role ("you are a senior tax attorney"). It can impose constraints ("do not invent information you are not given; if unsure, say so"). All of this measurably improves output quality. None of it adds new information.

Key insight

Prompt engineering can change *how* a model uses what it knows. It cannot change *what* it knows. If the answer to your question is not somewhere in the model's training weights, no amount of clever prompting will conjure it into existence.

Concretely, prompt engineering cannot solve the four hallucination examples we listed in section 1.4. It cannot produce the correct warranty period for your XR-440 controller, because the model has never seen your warranty document. It cannot quote a real Smith et al. 2023 paper, because the model has no way to fetch papers. You can add the constraint "only cite real sources," and the model may comply by refusing to cite anything — useful, but not the same as actually answering.

The instruction "answer based only on the provided context" is a piece of prompt engineering that becomes powerful precisely when context is provided to it. And providing relevant context, automatically and on demand, is what RAG does. So, prompt engineering and RAG are not competitors; they are partners. RAG supplies the knowledge; prompt engineering tells the model how to use it.

1.7 Why fine-tuning alone is not enough

If prompt engineering changes how the model uses its knowledge, fine-tuning changes the knowledge itself — at least in principle. By continuing the training process on a smaller, domain-specific dataset, you can

shift the model's weights so that it produces outputs more aligned with your target style, vocabulary, or domain.

Definition — Fine-tuning

The process of taking a pretrained LLM and training it further on a smaller, curated dataset to specialize its behavior — its tone, its format, its domain knowledge, or its alignment with a particular task. Fine-tuning updates the model's weights, in contrast to prompt engineering, which only changes the input.

Fine-tuning is genuinely useful for some problems — particularly tone, format, and consistent domain style. The mistake is treating it as the right tool for keeping facts up to date or for grounding answers in a specific corpus. For those jobs, it loses to RAG on three independent dimensions: cost, data, and frequency.

Cost

Fine-tuning is expensive. Even with parameter-efficient methods like LoRA, you need GPU compute, engineering time, and an evaluation pipeline to confirm the new model has not regressed on capabilities you care about. Full fine-tuning of a frontier-scale model is out of reach for most organizations entirely. RAG, by contrast, requires no model training — only an embedding pass over your documents, which can run on a laptop for thousands of pages.

Data

Effective fine-tuning needs a high-quality, curated training dataset — typically thousands of input-output pairs that exemplify exactly the behavior you want. Producing such a dataset is itself a significant project: collection, cleaning, labeling, deduplication, evaluation. Most teams that say "let's fine-tune" stall at this step. RAG bypasses the problem entirely: it works directly on your existing documents, in their existing form.

Frequency

This is the decisive one. Imagine your knowledge base is updated daily — new product releases, new policy memos, new support tickets. Fine-tuning every day is operationally impossible. Even fine-tuning weekly is a heavy lift. With RAG, an update is a file copy and a re-ingestion run. Minutes, not weeks.

Dimension	Fine-tuning	RAG
Changes model weights	Yes	No
Cost per knowledge update	High (GPU + eng time)	Low (re-embed new docs)
Data preparation effort	High (curated Q&A pairs)	Low (use existing docs as-is)
Update frequency feasible	Weekly to quarterly	Real-time to hourly
Source citations possible	No (knowledge is fused)	Yes (chunks are traceable)

Dimension	Fine-tuning	RAG
Best for	Tone, format, domain style	Facts, fresh data, citations
Hallucination reduction	Modest	Substantial (with grounding)

Table 1.1 — Fine-tuning vs RAG, head-to-head, for the job of grounding an LLM in a specific or evolving knowledge base.

The two techniques are not mutually exclusive — many production systems combine them, fine-tuning for tone and using RAG for facts. But for the specific problems of hallucination, knowledge cutoffs, and proprietary data, RAG is the strictly better tool. We will return to fine-tuning briefly in the final part of the book, when we discuss true weight-level self-improvement.

1.8 Where RAG fits in the landscape

We have now described every problem RAG was designed to solve. Let us assemble the picture and place RAG precisely where it belongs.

A plain LLM is a closed system: question goes in, answer comes out, and whatever knowledge the model has is whatever was baked into its weights at training time. Prompt engineering decorates the question more thoughtfully; fine-tuning rewrites parts of the model's internal weights. Neither lets the model consult anything new at the moment of answering. RAG does exactly that — and it does it without retraining anything.

Definition — Retrieval-Augmented Generation (RAG)

An architecture in which an LLM is paired with an external knowledge store. At query time, the system retrieves the most relevant pieces of information from the store and injects them into the model's context, so that the LLM generates an answer grounded in those specific, traceable sources rather than relying solely on its training weights.




Figure 1.2 — RAG at thirty thousand feet. The user's question is embedded, used to retrieve relevant chunks from a vector database, and the LLM answers using those chunks as evidence rather than relying solely on its training weights.

Read row by row, the diagram says: the user's question is converted into an embedding; that embedding is used to find the most semantically similar chunks in a vector database built from your own documents; those chunks are inserted into the prompt alongside the question; and the LLM generates its answer with that evidence sitting directly in front of it. Crucially, the model can be instructed to cite the chunks it used — which is why RAG answers come with traceable sources, while plain LLM answers do not.

What RAG fixes (and what it does not)

- **Hallucination** — significantly reduced. The model has real text in front of it; it no longer needs to invent.
- **Knowledge cutoffs** — neutralized. The vector store can be updated independently of the model.
- **Stale information** — neutralized. Re-ingest the new document, and the next query sees the new fact.
- **Proprietary or private data** — solved. The model never had to be trained on it; it only sees relevant chunks at query time.
- **Source citations** — solved. Every retrieved chunk has a known source, page, and metadata.

And what RAG does not fix:

- **Reasoning errors** — if the model is bad at multi-step logic, RAG cannot make it good.
- **Tone, style, or instructions** — those are prompt-engineering or fine-tuning concerns.
- **Bad retrieval** — if the vector search returns irrelevant chunks, the model is grounded in the wrong evidence and may answer worse than a plain LLM. Most of the rest of this book is about preventing that.

- **Aggregation and computation** — "What is the average age in this CSV?" is not a retrieval problem; it is a calculation problem. We will dedicate an entire chapter (the two-track architecture) to that distinction.

Where this book is taking you

By the end of this book you will have built — not just understood but built — a system that goes well beyond classical RAG. The agent will decide on its own when to retrieve, what to retrieve, and how many times. It will detect when its own answers are weak, learn from user feedback, remember which queries failed before, and distill validated interactions into long-term memory. RAG is the foundation. Everything else in this book stacks on top of it.

Chapter summary

LLMs are next-token predictors trained on enormous text corpora. They are powerful, fluent, and structurally limited in three ways relevant to production use: they hallucinate, they have a hard knowledge cutoff, and they cannot natively access proprietary data. Prompt engineering improves how a model uses what it knows but cannot supply new knowledge. Fine-tuning can update a model's weights but is too expensive, too data-hungry, and too slow for evolving facts. RAG closes exactly the gap that those tools cannot — by giving the model an external, queryable, updatable knowledge store, consulted on every question, with answers grounded in traceable sources.

In the next chapter, we will open up that abstract diagram and look at the actual moving parts: the ingestion pipeline, the retrieval pipeline, and the six core components — loader, splitter, embedder, vector DB, retriever, generator — that any RAG system, simple or sophisticated, is built from.